

Assignment 3

Varaprasad Kurra

Panther ID: 002430487

Sravanthi Malepati

Panther ID: 002538438

Olusesi Balogun


Panther ID: 002541505

Question 1:

Write a CUDA program for calculating the pre_x sums of an arbitrary-length array as described in Section 8.6 of the Kirk and Hwu book. You should follow their code snippets and input-output conventions.

Solution: In this problem, we have an arbitrary size array of length N. We have implemented our solution in 3 steps which includes Kogge – Stone implementation as one of its important steps. We will look into each step below

1. A total of 16 input elements are divided into four scan blocks.
2. We can use the Kogge–Stone kernel, the Brent–Kung kernel, or the three-phase kernel to process the individual scan blocks. The kernel treats the four scan blocks as independent input data sets.
3. After the scan kernel terminates operation, each Y element contains the scan result within its scan block.



An example specifying the illustration of given example.

Implementation:

We have defined the constant's in our program as follows

```
#define N_size 16          //number of elements in array
#define thread_number 4    //number of threads per block
#define block_number 4     //number of blocks
```

We have defined the following functions to implement the Algorithm

```
double myDiffTime(struct timeval &start, struct timeval &end)
{
    double d_start, d_end;
    d_start = (double)(start.tv_sec + start.tv_usec/1000000.0);
    d_end = (double)(end.tv_sec + end.tv_usec/1000000.0);
    return (d_end - d_start);
}

__global__ void prescan(float *gpu_outdata, float *gpu_indata, int n);
void scanCPU(float *f_out, float *f_in, int i_n);
```

Using the threads for carrying out the functionality.

```
for (int d = 1; d < thread_number; d *= 2)
{
    offset >>= 1;
    __syncthreads();
    if (thid < d)
    {
        int ai = bid * thread_number + offset*(2*thid+1)-1;
        int bi = bid * thread_number + offset*(2*thid+2)-1;
        float t = temp[bid * thread_number + ai];
        temp[ai] = temp[bi];
        temp[bi] += t;
    }
}
__syncthreads();
```

Scanning the array for carryout the Prefix Sums.

```
void scanCPU(float *f_out, float *f_in, int i_n)
{
    f_out[0] = 0;
    for (int i =1; i <=i_n; i++)
    {
        f_out[i] = f_out[i-1] + f_in[i-1];
    }
}
```

Finally Obtaining the output.

```
[smalepatil@cder01 ~]$ ./ques_1.exe
each element of an array a[0] = 2.000000
each element of an array a[1] = 1.000000
each element of an array a[2] = 3.000000
each element of an array a[3] = 1.000000
each element of an array a[4] = 0.000000
each element of an array a[5] = 4.000000
each element of an array a[6] = 1.000000
each element of an array a[7] = 2.000000
each element of an array a[8] = 0.000000
each element of an array a[9] = 3.000000
each element of an array a[10] = 1.000000
each element of an array a[11] = 2.000000
each element of an array a[12] = 5.000000
each element of an array a[13] = 3.000000
each element of an array a[14] = 1.000000
each element of an array a[15] = 2.000000
c[0] = 0.000, g[0] = 31.000
c[1] = 2.000, g[1] = 0.000
c[2] = 3.000, g[2] = 0.000
c[3] = 6.000, g[3] = 0.000
c[4] = 7.000, g[4] = 0.000
c[5] = 7.000, g[5] = 0.000
c[6] = 11.000, g[6] = 0.000
c[7] = 12.000, g[7] = 0.000
c[8] = 14.000, g[8] = 0.000
c[9] = 14.000, g[9] = 0.000
c[10] = 17.000, g[10] = 0.000
c[11] = 18.000, g[11] = 0.000
c[12] = 20.000, g[12] = 0.000
c[13] = 25.000, g[13] = 0.000
c[14] = 28.000, g[14] = 0.000
c[15] = 29.000, g[15] = 0.000
c[16] = 31.000, g[16] = 0.000
GPU Time for array size 16: 0.000078
CPU Time for array size 16: 0.000001
```

Question 2:

In this scan, a single kernel can be written to perform all three steps of the hierarchical scan algorithm. Thread block i first performs a scan on its scan block, using one of the three parallel algorithms of the above mentioned Question 1.

The block then waits for its left neighbor block $i-1$ to pass the sum value. Once the sum from block $i-1$ is received, the block generates and passes its sum value to its right neighbor block $i+1$.

The block then moves on to add the sum value received from block $i-1$ in order to complete all the output values of the scan block.

So the theme of this problem is to complete all the 3 steps performed by Kogge Stone Algorithm.

Code Snippet:

We have defined the following functions to execute the single pass scan.

```
__global__ void block_sum(const unsigned int * const gpu_in,unsigned int * const gpu_out,unsigned int * const gpu_sum,const size_t n);
__global__ void incrementeement( unsigned int * const d_array,const unsigned int * const d_incr);
```

Performing the single scan pass

```
__global__ void increment( unsigned int * const d_array,const unsigned int * const d_incr)
{
    const size_t bx = 2 * blockDim.x * blockIdx.x;
    const size_t tx = threadIdx.x;
    const unsigned int u = d_incr[blockIdx.x];
    d_array[bx + 2*tx] += u;
    d_array[bx + 2*tx+1] += u;
}
```

Main function to implement the Single pass scan

```
int main(int argc, char *argv[])
{
    const size_t len = 1000;
    thrust::host_vector<unsigned int> h_in(len);
    thrust::host_vector<unsigned int> h_out(len);
    for (size_t i = 0; i < h_in.size(); i++)
        h_in[i] = 3*i;
    psum(&h_in[0], &h_out[0], len);
    for (size_t i = 0; i < h_in.size(); i++)
        std::cout << h_in[i] << " " << h_out[i] << std::endl;
    return 0;
}
```